

Network Programming in Python


Mehmet Hadi Gunes




Modified from

- Steve Holden
- Dan Rubenstein
- DK Moon

- Review principles of networking
- Contrast TCP and UDP features
- Show how Python programs access networking functionality
- Give examples of client and server program structures
- Demonstrate some Python network libraries
- Give pointers to other network functionality



- Each layer uses the layer below
 - The lower layer adds headers to the data from the upper layer
 - The data from the upper layer can also be a header on data from the layer above ...





IP Characteristics

- Datagram-based
 - Connectionless
- Unreliable
 - Best efforts delivery
 - No delivery guarantees
- Logical (32-bit) addresses
 - Unrelated to physical addressing
 - Leading bits determine network membership




UDP Characteristics



- Also datagram-based
 - Connectionless, unreliable, can broadcast
- Applications usually message-based
 - No transport-layer retries
 - Applications handle (or ignore) errors
- Processes identified by port number
- Services live at specific ports
 - Usually below 1024, requiring privilege




TCP Characteristics



- Connection-oriented
 - Two endpoints of a virtual circuit
- Reliable
 - Application needs no error checking
- Stream-based
 - No predefined blocksize
- Processes identified by port numbers
- Services live at specific ports

- Various sockets... Any similarity?




- Endpoint of a connection
 - Identified by **IP address** and **Port number**
- Primitive to implement high-level networking interfaces
 - e.g., Remote procedure call (RPC)



Socket Programming API

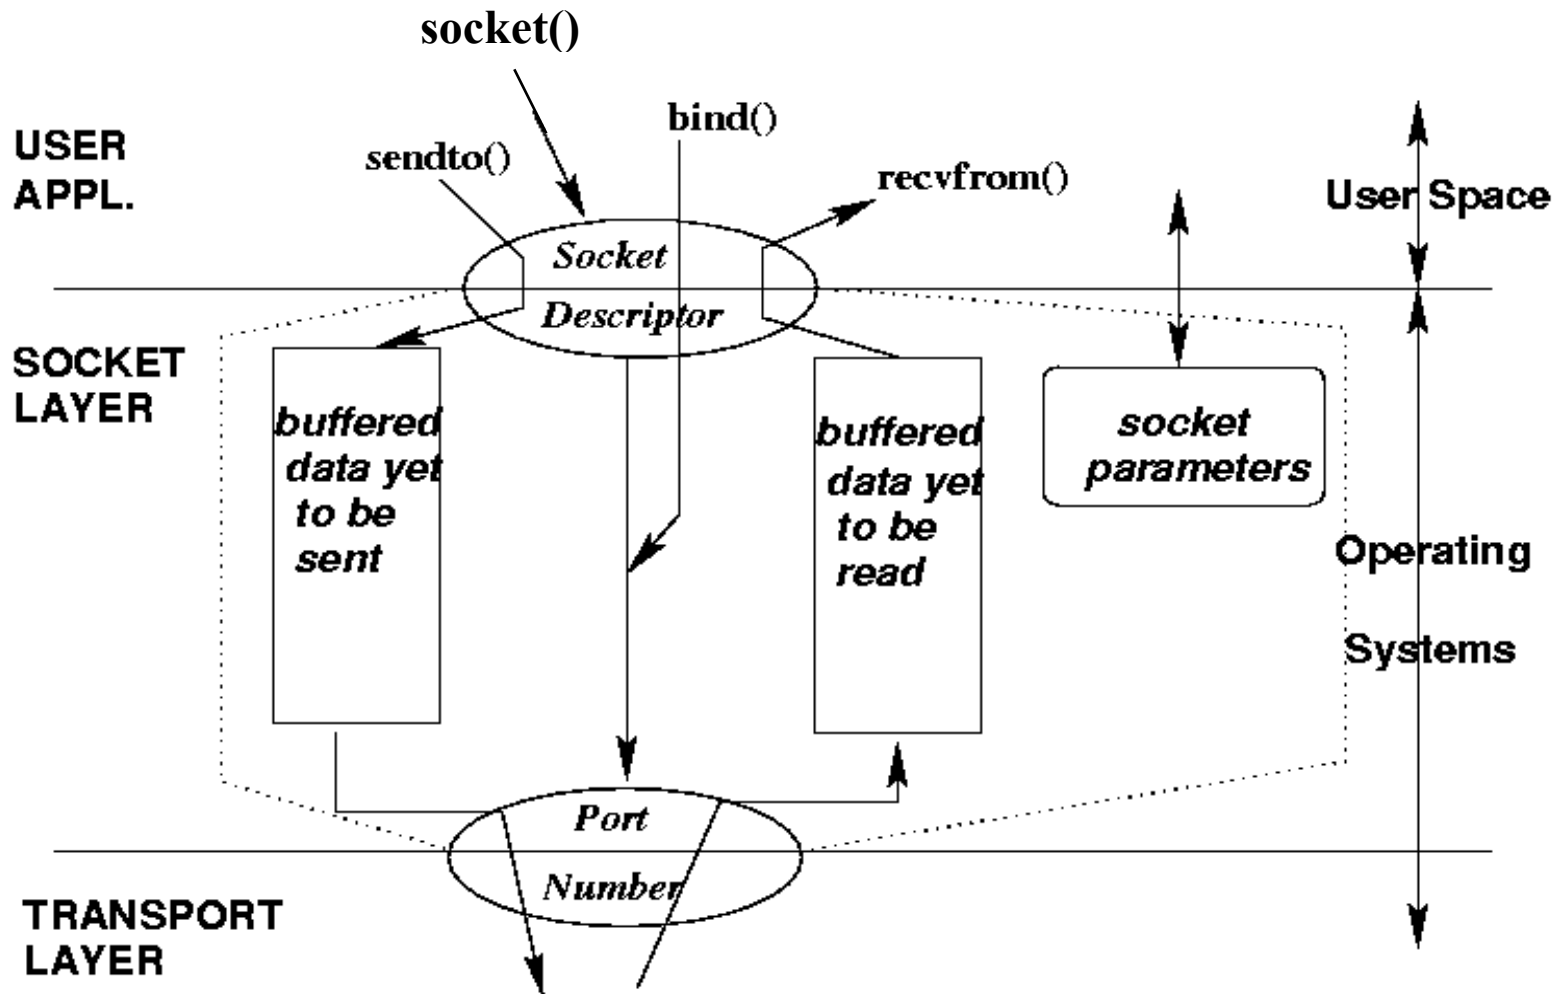
Application Programming Interface

- *Socket* analogous to door
 - sending process shoves message out door
 - sending process assumes transport infrastructure on other side of door which brings message to socket at receiving process
 - host-local, application created/owned, OS-controlled
 - connection between sockets set-up/managed by OS





Socket: Conceptual View




N

Two essential types of sockets


- SOCK_STREAM

- a.k.a. TCP
- reliable delivery
- in-order guaranteed
- connection-oriented
- bidirectional



- SOCK_DGRAM


- a.k.a. UDP
- unreliable delivery
- no order guarantees
- no notion of “connection” – app indicates dest. for each packet
- can send or receive




N

Types of Sockets

- When sending “Hi!” and “Hope you’re well”
- TCP treats them as a single bytes stream




- UDP treats them as separate messages



N

Types of Sockets (cont'd)

- Thus, TCP needs application-level message boundary.
 - By carrying length in application-level header





Client/Server Concepts

- Server opens a specific port
 - The one associated with its service
 - Then just waits for requests
 - Server is the passive opener
- Clients get ephemeral ports
 - Guaranteed unique, 1024 or greater
 - Uses them to communicate with server
 - Client is the active opener



A Socket-eye view of the Internet



cse.unr.edu
(134.197.20.22)



newworld.cs.umass.edu
(128.119.245.93)




cluster.cs.columbia.edu
(128.59.21.14, 128.59.16.7,
128.59.16.5, 128.59.16.4)

- Each host machine has an IP address
- When a packet arrives at a host

N

Ports

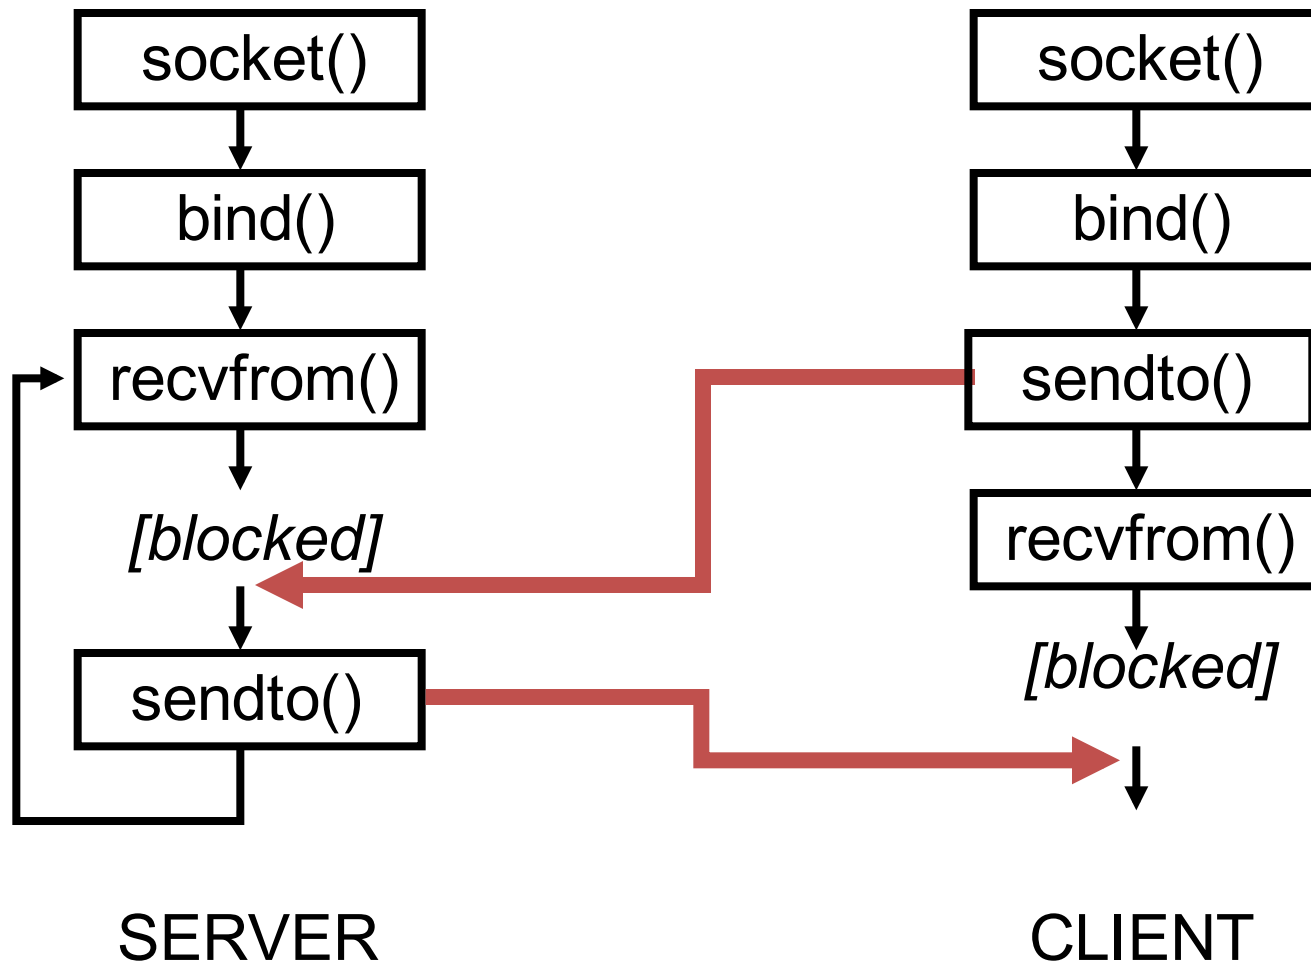
- Each host has 65,536 ports
- Some ports are *reserved for specific apps*
 - 20,21: FTP
 - 23: Telnet
 - 80: HTTP
 - see RFC 1700
 - about 2000 ports are reserved



A socket provides an interface to send data to/from the network through a port



Connectionless Services



N

Simple Connectionless Server

```
from socket import socket, AF_INET, SOCK_DGRAM
s = socket(AF_INET, SOCK_DGRAM)
s.bind(('127.0.0.1', 11111))
while True:
    data, addr = s.recvfrom(1024)
    print "Connection from", addr
    s.sendto(data.upper(), addr)
```

Empty -> all



Note that the *bind()* argument is a two-element tuple of address and port number

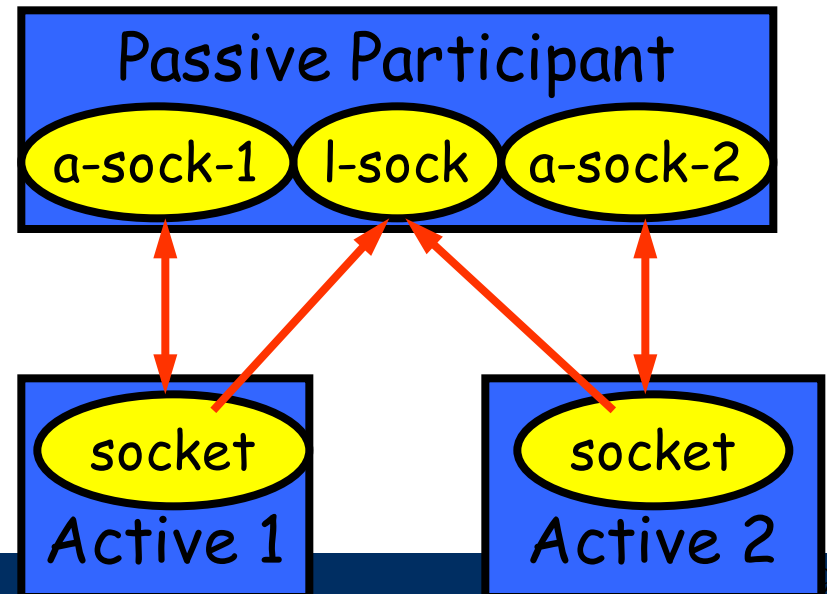


Simple Connectionless Client

```
from socket import socket, AF_INET, SOCK_DGRAM
s = socket(AF_INET, SOCK_DGRAM)
s.bind(('127.0.0.1', 0)) # OS chooses port
print "using", s.getsockname()
server = ('127.0.0.1', 11111)
s.sendto("MixedCaseString", server)
data, addr = s.recvfrom(1024)
print "received", data, "from", addr
s.close()
```

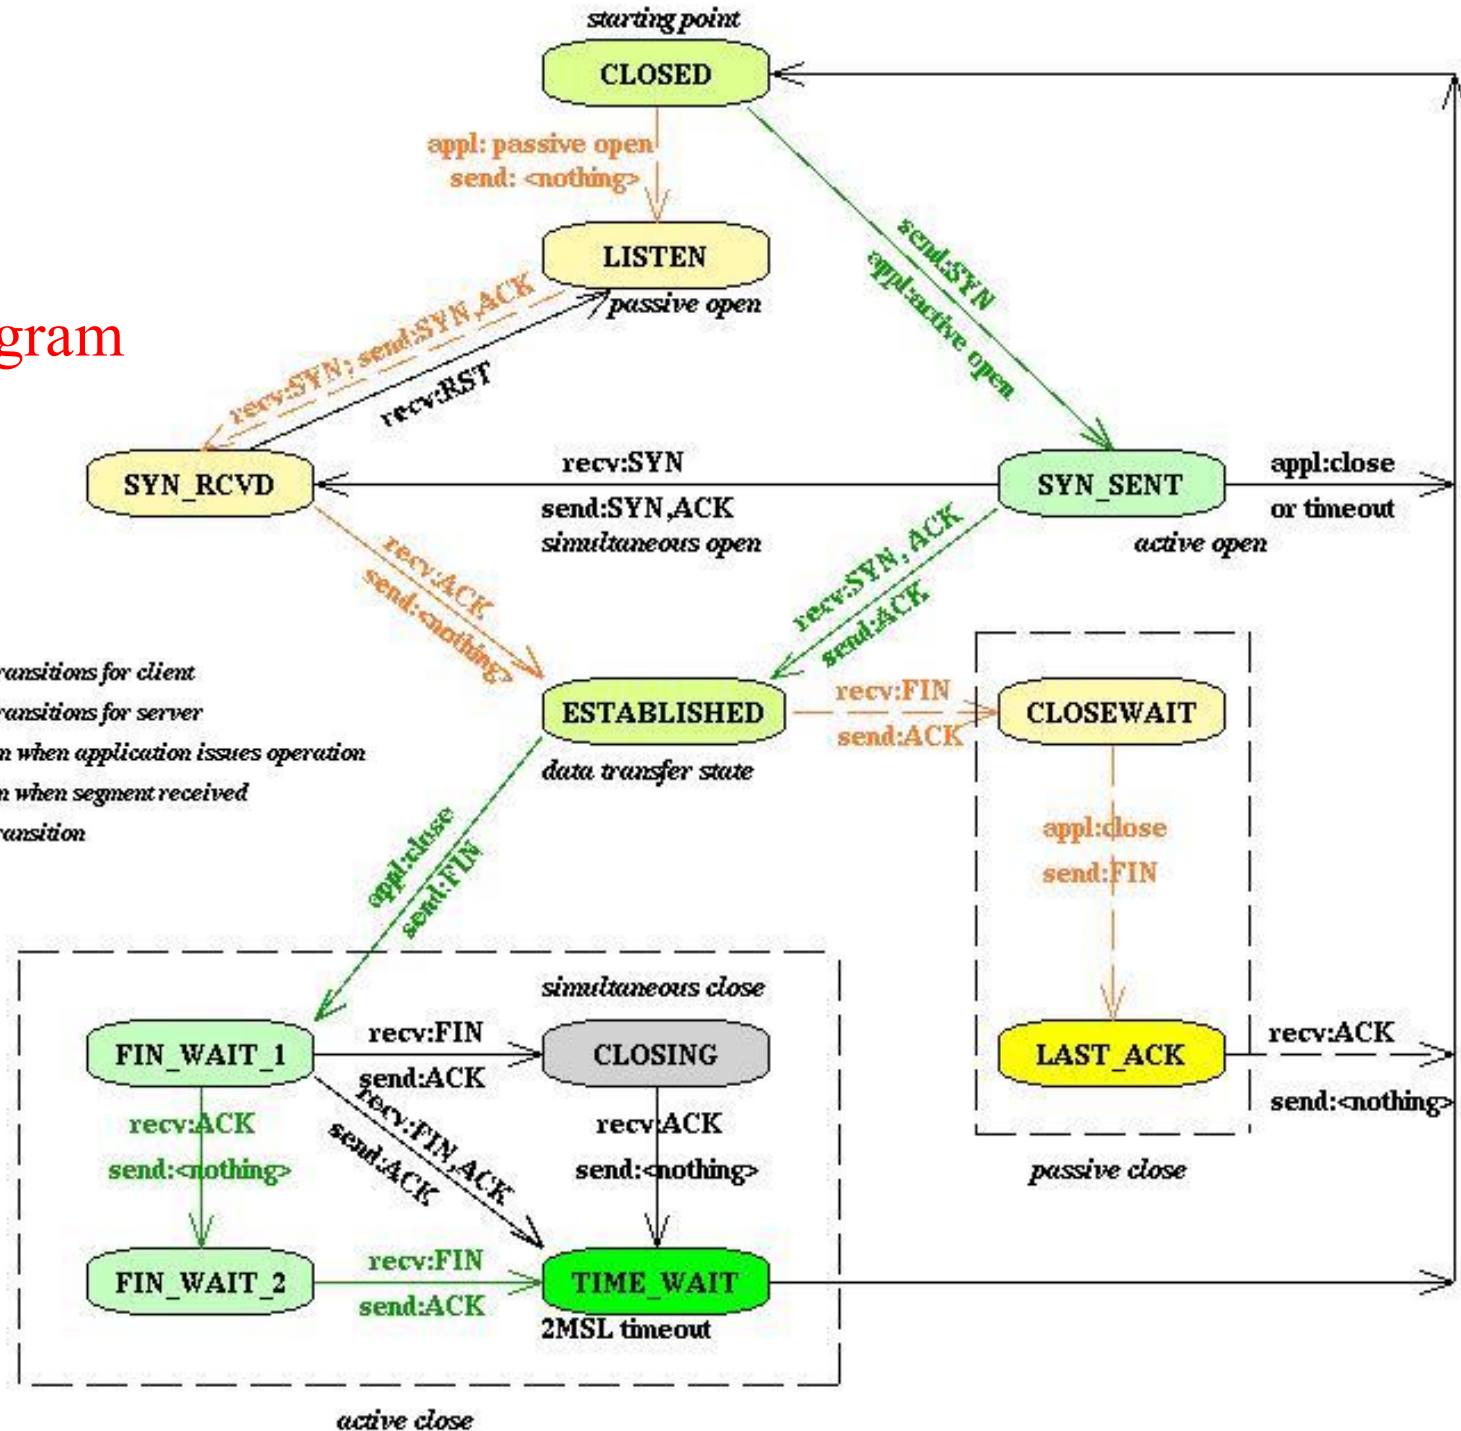
Connection setup cont'd

- Passive participant
 - step 1: **listen** (for incoming requests)
 - step 3: **accept** (a request)
 - step 4: data transfer
 - Active participant
 - step 2: request & establish **connection**
 - step 4: data transfer
- The accepted connection is on a new socket
- The old socket continues to listen for other active participants



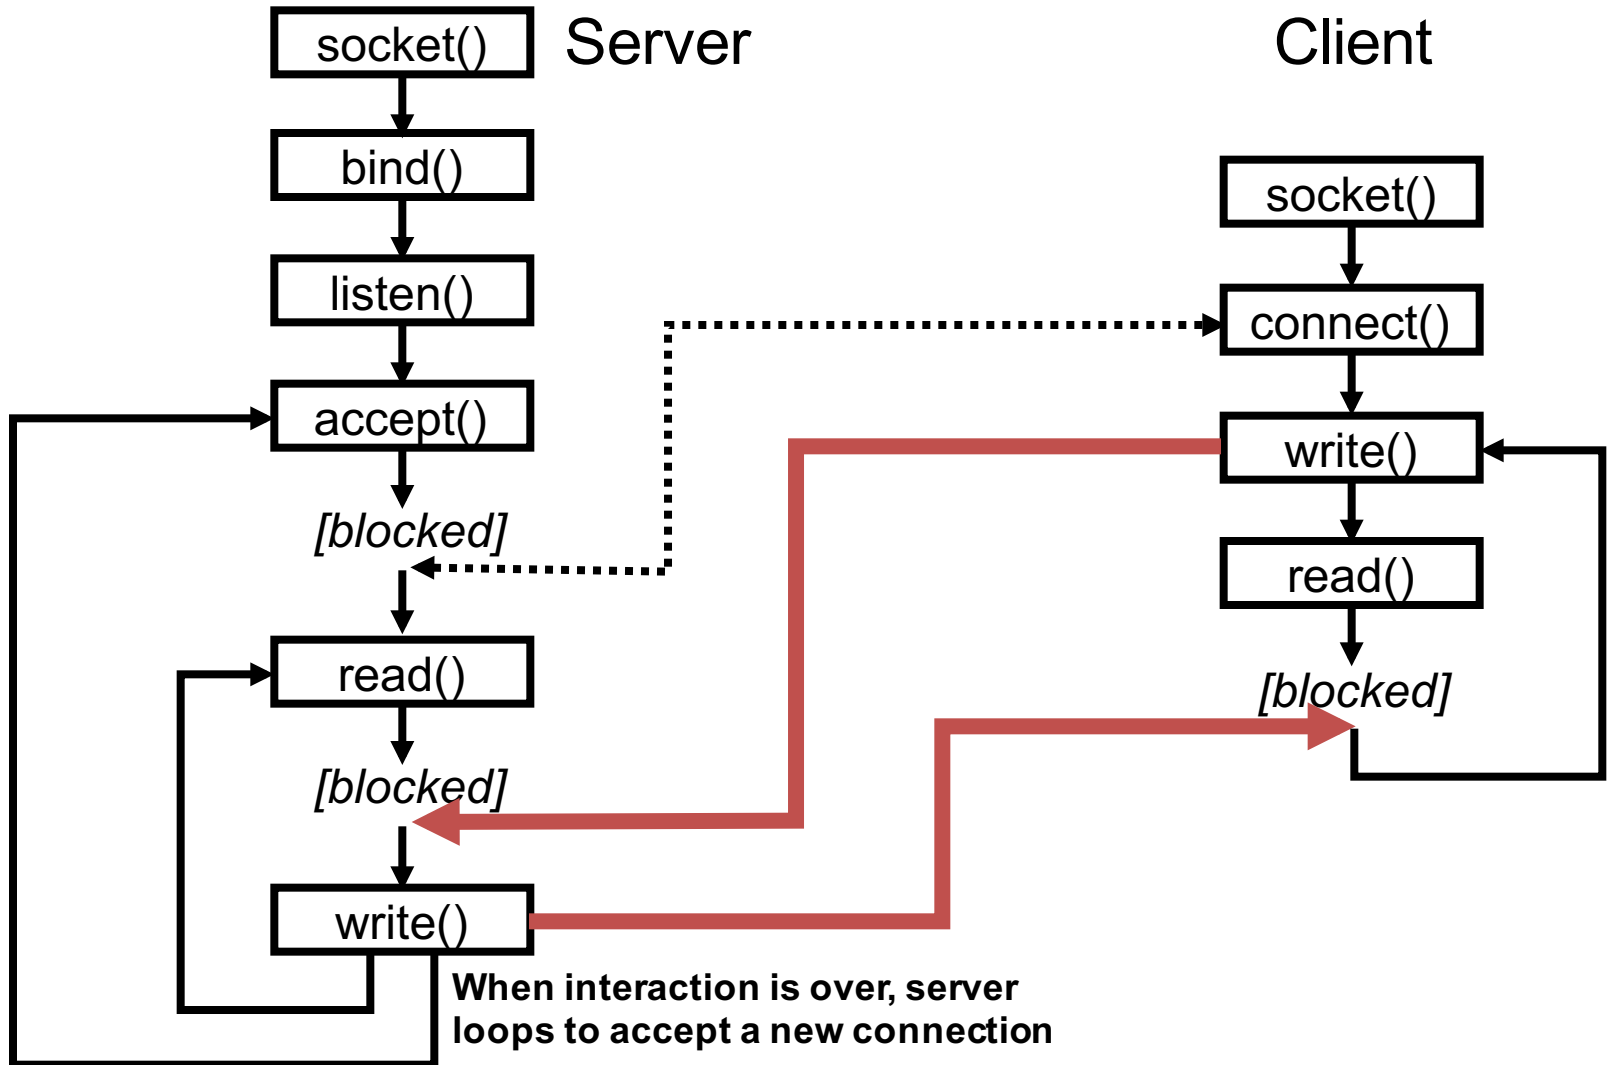


TCP state diagram





Connection-Oriented Services



```
from socket import \
    socket, AF_INET, SOCK_STREAM
s = socket(AF_INET, SOCK_STREAM)
s.bind(('127.0.0.1', 9999))
s.listen(5) # max queued connections
while True:
    sock, addr = s.accept()
    # use socket sock to communicate
    # with client process
```

- Client connection creates new socket
 - Returned with address by *accept()*
- Server handles one client at a time


```
s = socket(AF_INET, SOCK_STREAM)
s.connect((HOST, PORT))
s.send('Hello, world')
data = s.recv(1024)
s.close()
print 'Received', data
```

- This is a simple example
 - Sends message, receives response
 - Server receives 0 bytes after ***close()***

N

Some socket Utility Functions

- `htonl(i)` , `htons(i)`
 - 32-bit or 16-bit integer to network format
- `ntohl(i)` , `ntohs(i)`
 - 32-bit or 16-bit integer to host format
- `inet_aton(ipstr)` , `inet_ntoa(packed)`
 - Convert addresses between regular strings and 4-byte packed strings





Handling Names & Addresses

- `getfqdn(host='')`
 - Get canonical host name for host
- `gethostbyaddr(ipaddr)`
- `gethostbyname_ex(hostname)`
 - Returns (hostname, aliases, addresses)
 - Hostname is canonical name
 - Aliases is a list of other names
 - Addresses is a list of IP address strings



Treating Sockets as Files

- **`makefile([mode[, bufsize]])`**
 - Creates a file object that references the socket
 - Makes it easier to program to handle data streams
 - No need to assemble stream from buffers

- **socket.AF_UNIX**
 - Unix named pipe (NOT Windows...)
- **socket.AF_INET**
 - Internet – IP version 4
 - The basis of this class
- **socket.AF_INET6**
 - Internet – IP version 6
 - Rather more complicated ...

- **socket.SOCK_STREAM**
 - TCP, connection-oriented
- **socket.SOCK_DGRAM**
 - UDP, connectionless
- **socket.SOCK_RAW**
 - Gives access to subnetwork layer
- **SOCK_RDM, SOCK_SEQPACKET**
 - Very rarely used



Timeout Capabilities

- Can set a default for all sockets
 - `socket.setdefaulttimeout(seconds)`
 - Argument is float # of seconds
 - Or `None` (indicates no timeout)
- Can set a timeout on an existing socket `s`
 - `s.settimeout(seconds)`



Server Libraries

- `SocketServer` module provides basic server features
- Subclass the `TCPServer` and `UDPServer` classes to serve specific protocols
- Subclass `BaseRequestHandler`, overriding its `handle()` method, to handle requests
- Mix-in classes allow asynchronous handling via `ThreadingMixIn`

- Server instance created with address and handler-class as arguments:
`SocketServer.UDPServer(myaddr, MyHandler)`
- Each connection/transmission creates a request handler instance by calling the handler-class*
- Created handler instance handles a message (UDP) or a complete client session (TCP)

* In Python you instantiate a class by calling it like a function